

# Project BMOI: Morphogenic Osmotic Intelligence

Evolving Antifragile Cognitive Substrates via Neural Cellular Automata

Akshay John

Independent Researcher  
akshayjohn03@gmail.com

December 2025

## Abstract

Contemporary artificial intelligence systems rely on static architectures optimized via gradient descent. While effective, such systems are structurally brittle: partial damage often leads to catastrophic failure. Biological intelligence, by contrast, emerges from grown tissue and exhibits resilience, redundancy, and antifragility.

We introduce **Bio-Morphogenic Osmotic Intelligence (BMOI)**, a neural cellular automaton (NCA) substrate in which intelligence-like structure emerges as a stabilized morphogenic regime rather than an optimized function. Without backpropagation, labeled data, or weight transfer, the substrate evolves to internalize semantic geometry via manifold alignment from a frozen language model.

We empirically demonstrate three properties: scalability to  $120 \times 120$  grids, reproducible convergence across random seeds ( $\mu = 0.729, \sigma = 0.049$ ), and antifragility under extreme damage. BMOI constitutes a verified cognitive substrate upon which agent-level intelligence may later be constructed.

## 1 Introduction

Artificial intelligence research has largely equated progress with optimization: larger models, deeper networks, and more efficient training algorithms. However, these systems exhibit extreme fragility. Random deletion of a small fraction of parameters often causes total functional collapse, motivating the description of such models as *glass cannons*.

Biological cognition operates differently. Neural tissue is grown, not architected, and cognitive function persists despite injury and noise. This work explores whether intelligence-like structure can emerge from *morphogenesis*: local interaction rules that stabilize global function, rather than explicit optimization objectives.

We focus deliberately on the construction of a *cognitive substrate*, not an agent. The goal is not reasoning or planning, but the creation of a resilient medium capable of hosting such processes in the future.

## 2 Methodology

### 2.1 Neural Cellular Automaton Substrate

The system consists of a 2D grid of cells, each with a hidden state  $x_{ij} \in \mathbb{R}^d$ . Cell updates depend only on local information:

$$x_{t+1} = \tanh(\lambda x_t + \alpha \cdot \Delta(x_t))$$

where  $\lambda$  controls viscosity (memory decay),  $\alpha$  scales updates, and  $\Delta(\cdot)$  is a learned local reaction function. Updates are stochastic, introducing biological variability.

## 2.2 Semantic Osmosis

Rather than training on text, the organism is evolved to align its global activation vector with semantic target vectors derived from a frozen language model. No weights are copied. The substrate must rediscover semantic geometry through its own dynamics.

## 2.3 Fitness and Damage Protocols

Fitness is measured as cosine similarity between the organism’s global activation vector (mean grid state) and the semantic target.

Damage is implemented as stochastic cell death: at each timestep, individual cells are zeroed with probability  $p$ .

Training damage and evaluation damage are treated as independent experimental conditions.

# 3 Experimental Results

## 3.1 Phase 1: Scale Validation

To test scalability, the organism was deployed on a  $120 \times 120$  grid seeded with four independent semantic sources.




Figure 1: Planetary-scale test ( $120 \times 120$ ): four independent cognitive colonies persist without interference.

High viscosity created effective event horizons around each colony, preventing signal bleed at large spatial scales.

### 3.2 Phase 2: Reproducibility Benchmark

A five-seed benchmark assessed robustness. Semantic targets were fixed, while organism initialization and evolution varied.

Table 1: Genesis Benchmark ( $N = 5, 500$  Generations)

| Metric           | Mean  | Std Dev | Min   | Max   |
|------------------|-------|---------|-------|-------|
| Fitness (Cosine) | 0.729 | 0.049   | 0.664 | 0.798 |

### 3.3 Phase 3: Antifragility (Spartan Protocol)

A Naive organism (trained without damage) and a Spartan organism (trained under 20% chronic damage) were evaluated under increasing instantaneous damage.




Figure 2: Antifragility curve: performance of Naive vs Spartan organisms under increasing damage.

The Spartan organism exhibits hormesis: performance improves under moderate stress and remains high even at 90% damage.

## 4 Reproducibility Artifact

The following script generates the antifragility curve:

```

1 import pandas as pd
2 import matplotlib.pyplot as plt
3
4 df = pd.read_csv("bmoi_spartan_results.csv")
5
6 for label in df["Type"].unique():
7     subset = df[df["Type"] == label]
```

```

8     plt.plot(subset["Damage"], subset["Fitness"], marker='o', label=
9         label)
10 plt.xlabel("Damage Probability")
11 plt.ylabel("Fitness (Cosine Similarity)")
12 plt.title("Antifragility Curve: Naive vs Spartan")
13 plt.legend()
14 plt.grid(True)
15 plt.savefig("bmoi_spartan_curve.png", dpi=300)
16 plt.show()

```

Listing 1: Spartan Damage Curve Plot

## 5 Discussion and Future Work

BMOI establishes a cognitive substrate that is scalable, reproducible, and antifragile. It does not reason, plan, or act autonomously. Instead, it provides a resilient medium upon which higher-level intelligence may later be constructed.

Future work (Project Weaver) will explore transient binding mechanisms that allow spatially isolated cognitive regions to cooperate dynamically without collapsing modularity.

## 6 Conclusion

We demonstrate that intelligence-like structure can emerge as a stabilized morphogenic process rather than an optimized objective. BMOI grows, survives damage, and strengthens under stress—properties absent from contemporary neural architectures.

This work reframes AI research from designing minds to growing cognitive tissue.